

## Hierarchical (Agglomerative) Clustering - Vehicle Data Set using Scipy

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from scipy.cluster import hierarchy
from scipy.spatial import distance_matrix
from sklearn import manifold, datasets
from sklearn.cluster import AgglomerativeClustering
from scipy.cluster.hierarchy import fcluster

## Load data ##
df=pd.read_csv('D:\Python\edx\Machine Learning\Clustering\cars_clus.csv')
with open('hierarchical_vehicle.txt','a') as f:
    print(df.head(),file=f)
    print(df.shape,file=f)

## Data Cleaning ## clear the dataset by dropping the rows that have null value:

with open('hierarchical_vehicle.txt','a') as f:
    print('Shape of data set before cleaning: ',df.size,file=f)

df[[ 'sales', 'resale', 'type', 'price', 'engine_s',
     'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_cap',
     'mpg', 'lnsales']] = df[['sales', 'resale', 'type', 'price', 'engine_s',
     'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_cap',
     'mpg', 'lnsales']].apply(pd.to_numeric, errors='coerce')

df=df.dropna()
df=df.reset_index(drop=True)
with open('hierarchical_vehicle.txt','a') as f:
    print('Shape of the dataset after cleaning: ',df.size,file=f)
    print(df.head(5),file=f)

#Feature set
feat_set=df[['engine_s', 'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_c

# Normalization - between 0,1 for each feature using MinMaxScaler

from sklearn.preprocessing import MinMaxScaler
x=feat_set.values
min_max=MinMaxScaler()
feat_matrix=min_max.fit_transform(x)
```

```

with open ('hierarchical_vehicle.txt','a') as f:
    print(feat_matrix[0:5],file=f)

# First method - Clustering using Scipy
##In agglomerative clustering, at each iteration, the algorithm must update the
# distance matrix to reflect the distance of the newly formed cluster with the remaining
# clusters in the forest. The following methods are supported in Scipy for calculating
# the distance between the newly formed cluster and each:
# - single - complete - average - weighted - centroid

import scipy
leng=feat_matrix.shape[0]
D=scipy.zeros([leng,leng])
for i in range(leng):
    for j in range(leng):
        D[i,j]=scipy.spatial.distance.euclidean(feat_matrix[i],feat_matrix[j])

import pylab
Z=hierarchy.linkage(D,'complete')

# for partitioning in clustering we draw a cutting line
max_d=3
clusters=fcluster(Z,max_d,criterion='distance')
k=5
clusters_max=fcluster(Z,k,criterion='maxclust')
with open ('hierarchical_vehicle.txt','a') as f:
    print(clusters,file=f)
    print(clusters_max,file=f)

# Dendrogram
fig = pylab.figure(figsize=(18,50))
def llf(id):
    return ' [%s %s %s]' % (df['manufact'][id], df['model'][id], int(float(df['type'][id])))

dendro=hierarchy.dendrogram(Z,leaf_label_func=llf, leaf_rotation=0, leaf_font_size =4, o

#Display plot
plt.show()

```


Solution:

	manufact	model	sales	resale	type	price	engine_s	horsepow	wheelbas	width
length	curb_wgt	fuel_cap	mpg	lnsales	partition					
0	Acura	Integra	16.919	16.360	0.000	21.500	1.800	140.000	101.200	67.300
172.400	2.639	13.200	28.000	2.828		0.0				

```

1    Acura      TL  39.384  19.875  0.000  28.400    3.200  225.000  108.100  70.300
192.900  3.517  17.200  25.000  3.673    0.0
2    Acura      CL  14.114  18.225  0.000  $null$    3.200  225.000  106.900  70.600
192.000  3.470  17.200  26.000  2.647    0.0
3    Acura      RL   8.588  29.725  0.000  42.000    3.500  210.000  114.600  71.400
196.600  3.850  18.000  22.000  2.150    0.0
4    Audi       A4  20.397  22.255  0.000  23.990    1.800  150.000  102.600  68.200
178.000  2.998  16.400  27.000  3.015    0.0
(159, 16)
Shape of data set before cleaning: 2544
Shape of the dataset after cleaning: 1872
   manufact  model  sales  resale  type  price  engine_s  horsepower  wheelbas  width
length  curb_wgt  fuel_cap  mpg  lnsales  partition
0    Acura  Integra  16.919  16.360  0.0  21.50      1.8      140.0      101.2
67.3   172.4      2.639      13.2  28.0  2.828      0.0
1    Acura      TL  39.384  19.875  0.0  28.40      3.2      225.0      108.1
70.3   192.9      3.517      17.2  25.0  3.673      0.0
2    Acura      RL   8.588  29.725  0.0  42.00      3.5      210.0      114.6
71.4   196.6      3.850      18.0  22.0  2.150      0.0
3    Audi       A4  20.397  22.255  0.0  23.99      1.8      150.0      102.6
68.2   178.0      2.998      16.4  27.0  3.015      0.0
4    Audi       A6  18.780  23.555  0.0  33.95      2.8      200.0      108.7
76.1   192.0      3.561      18.5  22.0  2.933      0.0
[[0.11428571 0.21518987 0.18655098 0.28143713 0.30625832 0.2310559
  0.13364055 0.43333333]
 [0.31428571 0.43037975 0.3362256  0.46107784 0.5792277  0.50372671
  0.31797235 0.33333333]
 [0.35714286 0.39240506 0.47722343 0.52694611 0.62849534 0.60714286
  0.35483871 0.23333333]
 [0.11428571 0.24050633 0.21691974 0.33532934 0.38082557 0.34254658
  0.28110599 0.4      ]
 [0.25714286 0.36708861 0.34924078 0.80838323 0.56724368 0.5173913
  0.37788018 0.23333333]]
[ 1  5  5  6  5  4  6  5  5  5  5  4  4  5  1  6  5  5  5  4  2  11  6
  6  5  6  5  1  6  6  10  9  8  9  3  5  1  7  6  5  3  5  3  8  7  9  2
  6  6  5  4  2  1  6  5  2  7  5  5  5  4  4  3  2  6  6  5  7  4  7  6
  6  5  3  5  5  6  5  4  4  1  6  5  5  5  6  4  5  4  1  6  5  6  6  5
  5  5  7  7  7  2  2  1  2  6  5  1  1  1  7  8  1  1  6  1  1]
[1 3 3 3 3 2 3 3 3 3 3 2 2 3 1 3 3 3 3 2 1 5 3 3 3 3 3 1 3 3 4 4 4 4 2 3
 1 3 3 3 2 3 2 4 3 4 1 3 3 3 2 1 1 3 3 1 3 3 3 3 2 2 2 1 3 3 3 3 2 3 3 3 3
 2 3 3 3 3 2 2 1 3 3 3 3 3 2 3 2 1 3 3 3 3 3 3 3 3 3 3 3 1 1 1 1 3 3 1 1 1 3
 4 1 1 3 1 1]

```



Hierarchical (Agglomerative) Clustering - Vehicle Data Set using Scikit-learn

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from scipy.cluster import hierarchy
from scipy.spatial import distance_matrix
from sklearn import manifold, datasets
from sklearn.cluster import AgglomerativeClustering
from scipy.cluster.hierarchy import fcluster

## Load data ##
df=pd.read_csv('D:\Python\edx\Machine Learning\Clustering\cars_clus.csv')
with open('hierarchical_vehicle_s.txt','a') as f:
    print(df.head(),file=f)
    print(df.shape,file=f)

## Data Cleaning ## clear the dataset by dropping the rows that have null value:

with open('hierarchical_vehicle_s.txt','a') as f:
    print('Shape of data set before cleaning: ',df.size,file=f)

df[['sales', 'resale', 'type', 'price', 'engine_s',
```

```

        'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_cap',
        'mpg', 'lnsales']] = df[['sales', 'resale', 'type', 'price', 'engine_s',
        'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_cap',
        'mpg', 'lnsales']].apply(pd.to_numeric, errors='coerce')

df=df.dropna()
df=df.reset_index(drop=True)
with open ('hierarchical_vehicle_s.txt','a') as f:
    print('Shape of the dataset after cleaning: ',df.size,file=f)
    print(df.head(5),file=f)

#Feature set
feat_set=df[['engine_s', 'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_c

# Normalization - between 0,1 for each feature using MinMaxScaler

from sklearn.preprocessing import MinMaxScaler
x=feat_set.values
min_max=MinMaxScaler()
feat_matrix=min_max.fit_transform(x)
with open ('hierarchical_vehicle_s.txt','a') as f:
    print(feat_matrix[0:5],file=f)

# Second method - Clustering using scikit-learn
d_mat=distance_matrix(feat_matrix,feat_matrix)
with open ('hierarchical_vehicle_s.txt','a') as f:
    print(d_mat,file=f)

#AgglomerativeClustering performs a hierarchical clustering using a bottom up approach.
# The linkage criteria determines the metric used for the merge strategy:
# Ward minimizes the sum of squared differences within all clusters. It is a variance-mi
# approach and in this sense is similar to the k-means objective function but tackled wi
# agglomerative hierarchical approach.
# Maximum or complete linkage minimizes the maximum distance between observations of pai
# Average linkage minimizes the average of the distances between all observations of pai

agglom=AgglomerativeClustering(n_clusters=6,linkage='complete')
agglom.fit(feat_matrix)
with open ('hierarchical_vehicle_s.txt','a') as f:
    print(agglom.labels_,file=f)

# Adding new column - cluster to the data
df['cluster_']=agglom.labels_
with open ('hierarchical_vehicle_s.txt','a') as f:
    print(df.head(),file=f)

```

```

## Plotting scatter plot for data points with their clusters
import matplotlib.cm as cm
n_clusters=max(agglom.labels_)+1
colors=cm.rainbow(np.linspace(0,1,n_clusters))
cluster_labels=list(range(0,n_clusters))

plt.figure(figsize=(16,14))
for color, label in zip(colors,cluster_labels):
    subset=df[df.cluster_==label]
    for i in subset.index:
        plt.text(subset.horsepow[i], subset.mpg[i],str(subset['model'][i]), rotation=25)
    plt.scatter(subset.horsepow, subset.mpg, s= subset.price*10, c=color, label='cluster')
# plt.scatter(subset.horsepow, subset.mpg)
plt.legend()
plt.title('Clusters')
plt.xlabel('horsepow')
plt.ylabel('mpg')

# Centroids of each cluster is not clear in scatter plot, so we can summarize first
#classes and then the clusters. There are two classes - Cars and Trucks

qdf=df.groupby(['cluster_','type'])['cluster_'].count()
with open ('hierarchical_vehicle_s.txt','a') as f:
    print(qdf,file=f)

#For characteristics of each cluster
agg_cars=df.groupby(['cluster_','type'])['horsepow','engine_s','mpg','price'].mean()
with open ('hierarchical_vehicle_s.txt','a') as f:
    print(agg_cars,file=f)

##It is obvious that we have 3 main clusters with the majority of vehicles in those.
##Cars:
    ##Cluster 1: with almost high mpg, and low in horsepower.
    ##Cluster 2: with good mpg and horsepower, but higher price than average.
    ## Cluster 3: with low mpg, high horsepower, highest price.
##Trucks:
    ##Cluster 1: with almost highest mpg among trucks, and lowest in horsepower and price.
    ##Cluster 2: with almost low mpg and medium horsepower, but higher price than average.
    ##Cluster 3: with good mpg and horsepower, low price.

plt.figure(figsize=(16,10))
for color, label in zip(colors, cluster_labels):
    subset = agg_cars.loc[(label,)]
    for i in subset.index:
        plt.text(subset.loc[i][0]+5, subset.loc[i][2], 'type='+str(int(i)) + ', price='+
        plt.scatter(subset.horsepow, subset.mpg, s=subset.price*20, c=color, label='cluster')

```

```
plt.legend()
plt.title('Clusters')
plt.xlabel('horsepow')
plt.ylabel('mpg')

#Display plot
plt.show()
```

Solution:

	manufact	model	sales	resale	type	price	engine_s	horsepow	wheelbas	width
0	Acura	Integra	16.919	16.360	0.000	21.500	1.800	140.000	101.200	67.300
1	Acura	TL	39.384	19.875	0.000	28.400	3.200	225.000	108.100	70.300
2	Acura	CL	14.114	18.225	0.000	\$null\$	3.200	225.000	106.900	70.600
3	Acura	RL	8.588	29.725	0.000	42.000	3.500	210.000	114.600	71.400
4	Audi	A4	20.397	22.255	0.000	23.990	1.800	150.000	102.600	68.200

(159, 16)

Shape of data set before cleaning: 2544

Shape of the dataset after cleaning: 1872

	manufact	model	sales	resale	type	price	engine_s	horsepow	wheelbas	width
0	Acura	Integra	16.919	16.360	0.0	21.50	1.8	140.0	101.2	
1	Acura	TL	39.384	19.875	0.0	28.40	3.2	225.0	108.1	
2	Acura	RL	8.588	29.725	0.0	42.00	3.5	210.0	114.6	
3	Audi	A4	20.397	22.255	0.0	23.99	1.8	150.0	102.6	
4	Audi	A6	18.780	23.555	0.0	33.95	2.8	200.0	108.7	

```
[[0.11428571 0.21518987 0.18655098 0.28143713 0.30625832 0.2310559
 0.13364055 0.43333333]
 [0.31428571 0.43037975 0.3362256 0.46107784 0.5792277 0.50372671
 0.31797235 0.33333333]
 [0.35714286 0.39240506 0.47722343 0.52694611 0.62849534 0.60714286
 0.35483871 0.23333333]
 [0.11428571 0.24050633 0.21691974 0.33532934 0.38082557 0.34254658
 0.28110599 0.4 ]]
```

```

[0.25714286 0.36708861 0.34924078 0.80838323 0.56724368 0.5173913
 0.37788018 0.23333333]]
[[0.          0.57777143 0.75455727 ... 0.28530295 0.24917241 0.18879995]
 [0.57777143 0.          0.22798938 ... 0.36087756 0.66346677 0.62201282]
 [0.75455727 0.22798938 0.          ... 0.51727787 0.81786095 0.77930119]
 ...
 [0.28530295 0.36087756 0.51727787 ... 0.          0.41797928 0.35720492]
 [0.24917241 0.66346677 0.81786095 ... 0.41797928 0.          0.15212198]
 [0.18879995 0.62201282 0.77930119 ... 0.35720492 0.15212198 0.          ]]
[1 2 2 1 2 3 1 2 2 2 2 2 3 3 2 1 1 2 2 2 5 1 4 1 1 2 1 2 1 1 1 5 0 0 0 3 2
 1 2 1 2 3 2 3 0 3 0 1 1 1 2 3 1 1 1 2 1 1 2 2 2 3 3 3 1 1 1 2 1 2 2 1 1 2
 3 2 3 1 2 3 5 1 1 2 3 2 1 3 2 3 1 1 2 1 1 2 2 2 1 1 1 1 1 1 1 1 1 2 1 1 1 2
 0 1 1 1 1 1]
   manufact    model    sales    resale  type    price    engine_s    horsepower    wheelbas    width
length  curb_wgt  fuel_cap    mpg  lnsales  partition  cluster_
0      Acura  Integra  16.919   16.360   0.0    21.50         1.8         140.0         101.2
67.3    172.4      2.639     13.2   28.0    2.828         0.0          1
1      Acura      TL   39.384   19.875   0.0    28.40         3.2         225.0         108.1
70.3    192.9      3.517     17.2   25.0    3.673         0.0          2
2      Acura      RL   8.588   29.725   0.0    42.00         3.5         210.0         114.6
71.4    196.6      3.850     18.0   22.0    2.150         0.0          2
3       Audi      A4   20.397   22.255   0.0    23.99         1.8         150.0         102.6
68.2    178.0      2.998     16.4   27.0    3.015         0.0          1
4       Audi      A6   18.780   23.555   0.0    33.95         2.8         200.0         108.7
76.1    192.0      3.561     18.5   22.0    2.933         0.0          2
cluster_  type
0         1.0      6
1         0.0     47
          1.0      5
2         0.0     27
          1.0     11
3         0.0     10
          1.0      7
4         0.0      1
5         0.0      3
Name: cluster_, dtype: int64
           horsepower  engine_s          mpg          price
cluster_  type
0         1.0    211.666667   4.483333   16.166667   29.024667
1         0.0    146.531915   2.246809   27.021277   20.306128
          1.0    145.000000   2.580000   22.200000   17.009200
2         0.0    203.111111   3.303704   24.214815   27.750593
          1.0    182.090909   3.345455   20.181818   26.265364
3         0.0    256.500000   4.410000   21.500000   42.870400
          1.0    160.571429   3.071429   21.428571   21.527714
4         0.0     55.000000   1.000000   45.000000    9.235000

```



5	0.0	365.666667	6.233333	19.333333	66.010000
---	-----	------------	----------	-----------	-----------

